

Contents

I	Foundations	6
1	Introduction	8
2	Warm Up — Tic-Tac-Toe	9
2.1	Board Representation	9
2.2	Move representation	10
2.3	Search	10
2.4	Retrograde Analysis	10
3	Simple Player	11
3.1	Min-max search	11
3.2	Negamax search	14
3.3	Transposition tables	14
II	The Classical Model	15
4	Board representation, move generation	16
4.1	Piece lists	16
4.2	0x88	16
4.3	Mailbox	16
4.4	Bit-boards	16
4.4.1	Rotated bit-boards	16
4.4.2	Kindergarten bit-boards	16
4.4.3	Magic bit-boards	16
5	Search	17
5.1	Alpha-beta pruning	17
5.2	Principal variation	18
5.3	Transposition tables	18
5.3.1	Zobrist hashing	18
5.3.2	BCH hashing	18
5.4	Iterative deepening	18
5.5	Quiescence search	19
5.6	Null-move pruning	19
5.6.1	Adaptive null-move pruning	19
5.6.2	Verified null-move pruning	19
5.7	Other heuristics	19
5.7.1	Killer moves	19
5.7.2	History heuristic	19
5.7.3	Razoring	19
5.7.4	Futility pruning	19
5.7.5	Search extensions	19

6	Evaluation	20
6.1	Material	20
6.2	Piece-square bonus	20
6.3	Mobility	20
6.4	Structure, Shape	20
6.5	Square-control bonus	21
6.6	Hanging pieces, defended pieces	21
6.6.1	SEE	21
6.6.2	SOMA	21
6.6.3	SUPER-SOMA	21
III	The Monte Carlo Revolution	22
7	MCTS	23
7.1	UCB	23
8	Extensions	24
8.1	AMAF	24
8.2	RAVE	24
8.3	Incorporating domain knowledge	24
IV	Other	25
9	UI	26
9.1	GGS	26
9.2	UCI	26
9.3	UGI	26
9.4	UGE	26
9.5	XBoard	26
9.6	Fritz	26
10	Time management	27
10.1	Time controls	27
10.1.1	Arimaa Time Controls Specification	27
10.2	Estimation of remaining time	29
10.3	Control	29
10.4	Pondering	29
10.5	Example: Crafty	29
10.6	Example: Gerbil	29
10.7	Example: Glaurung	32
11	Advanced Search	34
11.1	Trappy min-max	34
11.2	Negascout	34
11.3	Null-window search	34
11.4	Aspiration search	34
11.5	PVS	34
11.6	MTD(f)	34
11.7	Conspiracy number	34
11.7.1	Alpha-beta conspiracy search	35
11.8	Proof number	36
11.8.1	Modifications and enhancements	41
11.9	Other heuristics	41

12 Advanced Evaluation	42
12.1 TDleaf(λ)	42
12.1.1 TD(μ)	43
12.2 Neural networks	43
12.3 Genetic algorithms	43
12.4 Pattern database / brains	43
13 GUI	44
14 Parallel Processing	45
14.1 Youngest Brothers Wait	45
15 Stuff	46
15.1 Notation	46
15.1.1 Forsythe-Edwards Notation (FEN)	46
15.1.2 GBR code	46
15.1.3 PGN	46
15.2 Typesetting	46
15.3 Ratings	46
15.3.1 Elo	46
15.3.2 Glicko	46
15.4 Computer analysis	46
15.4.1 Blunder check	46
15.4.2 Auto-annotation	46
15.4.3 Combination finding	46
15.4.4 Automated problem construction	46
16 Opening Book	47
16.1 Basic format	47
16.2 Basic usage	47
16.2.1 Hash approach	47
16.3 Automatic construction	47
16.4 Learning, updating	47
17 Endgame Tablebase	48
17.1 Building	48
17.1.1 Thompson approach	48
17.1.2 Wu-Beal approach	49
17.2 Encoding: symmetries and numbering	50
17.3 Compressing	50
17.4 Using	51
V Specific Games	52
18 Standard Models	53
18.1 $\alpha\beta$ -search	53
18.2 MCTS	53
18.3 Etc.	53
19 Congo	54
19.1 Endgame Tablebase	54
19.1.1 Another attempt at counting endgame positions (10-20-2011)	55

20 Mancala	56
20.1 Kalah	56
20.2 Awari	56
20.3 Mancala	56
20.4 Dakon	56
20.5 Bao	56
21 Chess	57
22 Checkers	58
23 Xiangqi	59
24 Shōgi	60
24.1 Shōgi	60
24.2 Dōbutsu shōgi	60
24.3 Tori shōgi	60
24.4 Poppy shōgi	60
24.5 Whale shōgi	60
24.6 Chu shōgi	60
24.7 Tenjiku shōgi	60
25 Tafl	61
25.1 Evaluation	61
25.2 Mate search	61
26 Go	62
27 Havannah	63
28 Hex	64
29 Hive	65
30 Amazons	66
31 Dots-and-boxes	67
32 Exotica	68
32.1 Backgammon	68
32.2 Poker	68
32.3 Kriegspiel, Stratego	68
32.4 Wargames, Advanced Squad Leader	68
33 Combinatorial game theory	69
34 Zillions of Games	70
35 International Computer Olympiada (ICGA)	71
35.1 2008 Computer Olympiad (call for participation)	72
35.2 2009 Computer Olympiad (competed)	72
35.3 2010 Computer Olympiad	72
35.3.1 call for participation	72
35.3.2 participated	72

36 Game rules, sample games, comparisons	73
36.1 Congo	73
36.2 Chess	73
36.3 Shogi	74
Glossary	75
Index	76
Bibliography	78

Part I

Foundations

Need:

consistent notation for moves / trees / ...

vocabulary for move / board / game / state / evaluation / etc.

vocabulary for pieces / colors / etc.

build my own game-board typesetting macros?

pseudo-code / program-listing style...

do a couple full game rules

(fairy) chess boards

go boards

hex boards

C++ / Python / Scala / OCaml / C ???

USE TikZ package for graphs/trees?

Chapter 1

Introduction

We'll use C or C++ for programming language probably. Except if we don't.

Chapter 2

Warm Up — Tic-Tac-Toe

The simple game of TIC-TAC-TOE (or NOUGHTS-AND-CROSSES or ...) is well-known to most people and viewed as a simple child's game. Despite this, we'll ...

we will build various pieces for this

The simplicity of this game will allow us to look at a variety of aspects game programming without getting bogged down in complicated and game-specific details of rules or game representations. And, as it is such a simple game, we can do a deeper analysis of the game. There are less than $3^9 = 19683$ different possible board states the game can be in — each of the 9 squares can be empty, marked “X”, or marked “O” (three possibilities). The player to move can be determined by the number of X's and O's on the board (if they are the same, it is the first player's move, if different then the second player's.) Note that this count is an overestimate, since some of the counted possibilities are impossible in a game (for example, a board filled completely with X's.) In fact, the exact number of legal board positions that can occur in a game is ???TODO??? — we will compute this number in our work below).

2.1 Board Representation

The first step we'll need is to find a way to represent the state of the game. The most obvious approach is to simply make a three-by-three array to hold the state of each cell in the tic-tac-toe grid.

```
typedef enum piece { empty, x_piece, o_piece } piece;
typedef piece board[3][3];
```

A slight variation, which is sometimes convenient, is to collapse the two-dimensional array to a one-dimensional array:

```
typedef enum piece { empty, x_piece, o_piece } piece;
typedef piece board[9];
```

This means that if we want to use a row-column pair to access the cells, then we'll have to do the conversion ourself. We might define convenience functions or macros for this.

```
#define RC_to_I(r,c) ((r)*3+(c))
#define I_to_R(i)    ((i)/3)
#define I_to_C(i)    ((i)%3)
```

Though with

A less obvious approach, but taking much less space and allowing us to do some cute tricks in implemetation is a bit-board approach. That is, we use the individual bits in an integer to ... Note that we need to use an integer type which has at least 9 bits — one for each cell of the board. Note that this takes 4 bytes on a typical machine, whereas the cell approach takes 9 bytes.

```
typedef unsigned short bitboard;
typedef struct board {
    bitboard x;
    bitboard o;
} board;
```

2.2 Move representation

For tic-tac-toe, a move can simply be represented

2.3 Search

The first programming we'll do is write a solver for the game. This will tell us, for any position, the *theoretical value* of the position — that is, if we assume two players who play the best possible move at each position, whether the game is a win for the first player or a win for the second player or a draw. It is a theoretical fact, first proven by von Neumann, that one of these three possibilities must hold. Our algorithm will in fact be a kind of constructive “proof” of this.

2.4 Retrograde Analysis

An alternative to searching from a given position is to work backwards from the end of the game: find the positions that are already ended (win, loss, or draw). Then find positions that are one move away from ending; then two moves away; etc.

The basic technique is to list every possible position, mark the positions which are wins for X or O (and the invalid one) and then work backwards from those.

Chapter 3

Simple Player


In this chapter we build simple players for the games of MANCALA, MINI-CHESS, and MINI-SHOGI. This will show us the issues that can arise, but our implementation will be simplistic and naïve. Later chapters will extend and enhance ...

Our core algorithms will be general and will be reused for each of the games, with only game-specific modules needing to be changed for move-generation and position evaluation.

We'll need to

3.1 Min-max search

The basis of most game-playing programs is the “min-max” search which is based on the idea of searching exhaustively through every possible sequence of moves to the end of the game and ...



```
def Search(b: Board, depth: Int) {  
  if Terminal(b) return Result(b)  
  if (depth==0) return Eval(b)  
  if (b.ToMove == Maximizer) {  
    best_result = -Inf  
    for (m <- Moves(b)) {  
      MakeMove(b, m)  
      result = Search(b, depth-1)  
      UnMakeMove(b,m)  
      if (result > best_result) best_result = result;  
    }  
  } else { // b.ToMove == Minimizer  
    ...  
  }  
}  
  
def Search(g: Game, depth: Int, eval: Evaluator): Value = {  
  if (g.board.Terminal()) return g.board.Result()  
  ...  
}
```

```

if (depth==0) return eval(g.board)
g.board.ToMove() match {
  case Maximizer => {
    var best_v = Value.Min_Inf
    for (m <- g.board.Moves()) {
      g.board.MakeMove(m)
      var res = Search(g, depth-1, eval)
      g.board.UnMakeMove(m)
      res = res.step
      if (res > best_v) best_v = res
    }
    return best_v
  }
  case Minimizer => {
    var best_v = Value.Pos_Inf
    for (m <- g.board.Moves()) {
      g.board.MakeMove(m)
      var res = Search(g, depth-1, eval)
      g.board.UnMakeMove(m)
      res = res.step
      if (res < best_v) best_v = res
    }
    return best_v
  }
}
}

template <class TMove, bool tt>
result minimax(Board<TMove>& b) {
  hash h;
  if (tt) {
    h = b.Hash();
    if (ttable[h].type == EXACT_BD) {
      ++tthits;
      return ttable[h].res;
    }
  }
  result eval = b.Result();
  ++nodes_searched;
  if (IS_TERMINAL(eval)) {
    if (tt) { ttable[h].res = eval; ttable[h].type = EXACT_BD; }
    return eval;
  }
  result best_v;
  if (b.ToMove() == XX) {
    std::vector<TMove> moves;
    best_v = OO_WIN_IN(-1);
    int num_m = b.Moves(moves);
    for (int m=0; m<num_m; ++m) {
      b.MakeMove(moves[m]);
      result res = minimax<TMove,tt>(b);
      b.UnMakeMove(moves[m]);
      res = PLY_UP_RESULT(res);
      if (res > best_v) best_v = res;
    }
  }
}

```

```

    }
} else {
    std::vector<TMove> moves;
    best_v = XX_WIN_IN(-1);
    int num_m = b.Moves(moves);
    for (int m=0; m<num_m; ++m) {
        b.MakeMove(moves[m]);
        result res = minimax<TMove,tt>(b);
        b.UnMakeMove(moves[m]);
        res = PLY_UP_RESULT(res);
        if (res < best_v) best_v = res;
    }
}
if (tt) { ttable[h].res = best_v; ttable[h].type = EXACT_BD; }
return best_v;
}

template <class Board, class TMove, bool tt>
result minimax(Board& b) {
    hash h;
    if (tt) {
        h = b.Hash();
        if (ttable[h].type == EXACT_BD) {
            ++tthits;
            return ttable[h].res;
        }
    }
    result eval = b.Result();
    ++nodes_searched;
    if (IS_TERMINAL(eval)) {
        if (tt) { ttable[h].res = eval; ttable[h].type = EXACT_BD; }
        return eval;
    }
    result best_v;
    if (b.ToMove() == XX) {
        std::vector<TMove> moves;
        best_v = OO_WIN_IN(-1);
        int num_m = b.Moves(moves);
        for (int m=0; m<num_m; ++m) {
            b.MakeMove(moves[m]);
            result res = minimax<Board,TMove,tt>(b);
            b.UnMakeMove(moves[m]);
            res = PLY_UP_RESULT(res);
            if (res > best_v) best_v = res;
        }
    } else {
        std::vector<TMove> moves;
        best_v = XX_WIN_IN(-1);
        int num_m = b.Moves(moves);
        for (int m=0; m<num_m; ++m) {
            b.MakeMove(moves[m]);
            result res = minimax<Board,TMove,tt>(b);
            b.UnMakeMove(moves[m]);
            res = PLY_UP_RESULT(res);

```

```

    if (res < best_v) best_v = res;
  }
}
if (tt) { ttable[h].res = best_v; ttable[h].type = EXACT_BD; }
return best_v;
}


```

3.2 Negamax search

Most real game-playing programs don't use the min-max formulation for their search, but rather use the so-called "negamax" form. This gives a function which is equivalent (meaning it will search exactly the same nodes and will return exactly the same result) but requires less code. Instead of requiring two different functions (which are nearly identical), one for the Min player and one for the Max player, negamax search uses a single search routine which is always trying to maximize the outcome from the point of view of the player whose turn it is (rather than only player consistently trying to maximize and the other player trying to minimize).

Another way to think about it is that min-max search uses an absolute evaluation whereas negamax uses a relative evaluation: evaluating the position from the point-of-view of the person to move. The relative evaluation is closer to how we might actually think about the position if we were playing a game: we will think that it is good for us or bad for us. The absolute evaluation is locked to thinking in terms of a single player.

Generally negamax formulations, though sometimes confusing when first getting into this area, tend to lead to cleaner and more natural code. It eliminates a lot of if-then checks for which side is moving in the search code. (It can add a tiny complication to the board evaluation which is often naturally coded in absolute terms.)



3.3 Transposition tables

We can reduce significantly the amount of effort spent in min-max style searching by recognizing that many positions will arise multiple times in the course of our search — for example, one might move a piece twice over the course of two turns, ending up at the same position, but with a different intermediate step.

Part II

The Classical Model

Chapter 4

Board representation, move generation

Here

4.1 Piece lists

Piece lists are useful for larger games where the board density (ratio of number of pieces to squares on the board) is low. Instead of directly representing the board, we simply record for each piece, where it is on the board. It is often convenient to also represent the board explicitly, but can still use piece-lists in addition to speed up move generation or other calculations.

The advantage here is that when doing move generation we don't have to scan the entire board looking for pieces to move — we've got the pieces directly accessible and ready to iterate over. If we organize the lists correctly, we can also get the advantage of generating moves for certain kinds of pieces before or after other kinds of pieces.

4.2 0x88

This is a trick which works for 8×8 boards but does not nicely generalize to other sizes of boards (or to non-rectangular boards). It is fairly popular in chess playing programs. The idea is to note that representing a number between 0 and 7 requires exactly 3 bits and that when generating moves for chess, any putative moves that would take a piece off of the board would take them to a coordinate which sets the fourth (or “8”) bit. Thus, if we represent an x-y coordinate in bits as “0rrr0ccc” then we can detect an illegal coordinate simply by xor'ing with $0x88 = b10001000$.

4.3 Mailbox

4.4 Bit-boards

Bit-boards are a popular approach which are designed to take advantage of the fact that bit-wise operations are performed quickly and in parallel by the computer. The match of the 64 squares of the chess-board, for example, match very nicely with the 64 bits in a machine word in a modern 64-bit processor.

The basic idea is that for each piece type, we maintain a vector of bits, one bit for each square on the board, indicating if there is a piece of that type on the board.

4.4.1 Rotated bit-boards

4.4.2 Kindergarten bit-boards

4.4.3 Magic bit-boards

Chapter 5


Search

In this chapter we investigate more deeply the issue of searching,

5.1 Alpha-beta pruning

Alpha-beta ($\alpha\beta$) pruning is a way to get the same result as min-max (or negamax) searching but with significantly reduced amount of search. The idea is that we don't really need to search every node, but can ignore significant parts of the game-tree without affecting the accuracy of our result.


The basic idea is illustrated in the following tree:




where we can ignore the (subtree off the) final branch by the following reasoning: the max player can guarantee a result of *at least 2* by making the first move, while if he makes the second move, the min player can ensure a result of *no better than 1* by taking her first move, irrespective of the value of her second move. Thus the max player would not take his second move and we do not need to search any further.

The effectiveness of $\alpha\beta$ -pruning is extremely sensitive to the order in which we search the different possibilities at a given node. The best case is when the first node we search at each node is the best possible move from that position. Unfortunately, we typically do not know the best move (if we know the best move, there would be no need to search in the first place!)

Worst-case scenario (9 leaf nodes evaluated):



Best-case scenario (5 leaf nodes evaluated):



The theoretical improvement we can expect from full width search to optimal $\alpha\beta$ -search was computed by Knuth and [...] which takes a tree with uniform branching factor b of depth L from b^L nodes for complete search to $b^{\lceil L/2 \rceil} + b^{\lfloor L/2 \rfloor}$ which translates roughly into allowing us to search to *twice* the depth for the same amount of time!

5.2 Principal variation

The *principal variant* or *PV* is the expected optimal line of play, where both players are ...

5.3 Transposition tables

Transposition tables are critical for efficient performance.

A crucial element of using transposition tables is the hash function — the function which takes a position and

5.3.1 Zobrist hashing

The standard approach to hashing was ... by Zobrist in ... It is a clever approach which applies easily to many different games and implementation styles.

The core idea is to assign a (pseudo-)random number to each individual color, piece, and square combination. For example, in chess, we'll have a number (code) for white knight at a1, white knight at a2, ..., black rook at h8. Then, to compute the hash of a position, we simply take the “exclusive or” (or “xor” or \oplus) of this code for each piece on the board.¹

Now, the xor operator has the beautiful property that $(a \oplus b) \oplus b = a$ for any a and b . In other words, if I take a number and xor any other number with it twice, I end up with my original number. Additionally, xor operates bit-by-bit on an integer, so it is a very fast operation to execute by the CPU.

5.3.2 BCH hashing

5.4 Iterative deepening

In iterative-deepening, we do a series of searches to increasing depth.

As we will see later, using iterative deepening makes for good use of time — allowing us to do the best search we can in a given time interval, even though ...

Although re-searching the same position multiple times may seem like a ridiculous waste of time, it is actually quite efficient and it is possible for an iterative-deepened search to a given depth to actually take *less* time than a single search directly to that depth! The main reason for this is that with the use of a transposition table, we are able to use the results of shallower searches to sort the moves when doing a deeper search, thus tending to get better results for $\alpha\beta$ -pruning. Another point to note is that often shallower search results stored in the transposition table are re-used in the deeper search allowing for savings in the deeper search.

¹The xor of individual bits is defined by $0 \oplus 1 = 1 \oplus 0 = 1$ and $0 \oplus 0 = 1 \oplus 1 = 0$. The xor of two n -bit integers is simply the bit-by-bit xor of the individual bits in the integer representations. You may check that, for example, $1 \oplus 3 = 2$ and $1 \oplus 4 = 5$.

5.5 Quiescence search

The idea behind quiescence search is to end search lines only at “quiet” positions. This stabilizes evaluations and helps to avoid the so-called “horizon effect”. It will also add to the tactical strength.

As an example of the horizon effect, consider the following chess position:

...

If we use a standard fixed-depth alpha-beta search we get the following lines ... notice that it isn’t until we reach a search depth of ... that the search finally realizes that snagging the pawn isn’t safe. On the other hand using quiescence extensions we get ...

The definition of quiet or quiescent position depends on the game and the nature of the search can vary.

For CHESS, a common approach to quiescence search is to search all captures, checks, and check responses. Note that it is important to allow a “stand pat” option also when searching (where a side can just accept the static evaluation instead of following further in the quiescence search).

Following captures is a relatively safe and self-

For SHOGI, doing a quiescence search which searched captures as well as drops would lead to horrendous search explosions as capture-drop sequences can be unending.

5.6 Null-move pruning

Null-move pruning is a forward-pruning technique based on the assumption that it is always better to make a move than to “pass”. It can significantly reduce the number of nodes searched with a minor reduction in the strength (due to the occasional ...)

5.6.1 Adaptive null-move pruning

For chess it was ... $R = 3$ was somewhat over-aggressive in pruning, while $R = 2$ is

5.6.2 Verified null-move pruning

5.7 Other heuristics

Here we cover some heuristic pruning techniques which are “unsound” — unlike $\alpha\beta$ -pruning, they are not guaranteed to return the exact min-max value. However, in practice they are extremely effective.

5.7.1 Killer moves

5.7.2 History heuristic

5.7.3 Razoring

5.7.4 Futility pruning

Extended futility pruning

5.7.5 Search extensions

Chapter 6

Evaluation

In this chapter we discuss the evaluation of

6.1 Material

The most significant of components for evaluation of positions for most chess-like games is material: the positions ...

6.2 Piece-square bonus

It is known that in CHESS, for example, it is generally advantageous for knights to occupy central squares and that usually it is bad to have a knight in a corner square. This can be generalized to giving a bonus or malus for each piece being in any particular position on the board.

Similarly, certain general tactics can be “encouraged” by giving bonuses and maluses:

- General piece centralization
- Pawn pushing towards promotion squares
- Bishops on long diagonals (fianchetto)
- ...

6.3 Mobility

Mobility describes the available moves by each side. The simplest approach would simply count the possible moves in the given position. More sophisticated measures include counting attacks differently than simple moves, “x-ray” moves (moves or attacks which are blocked by one’s own pieces but which ... TODO...

6.4 Structure, Shape

For games such as CHESS or SHŌGI, we have “static” features of the position — features of the position which tend to persist over longer stretches of the game, features which involve the interaction of many pieces, ...

For CHESS, this would include pawn-structure (e.g. isolated pawns, double pawns, passed pawns, protected passed pawns, etc.) whereas

For SHOGI we have “shape” TODO...

6.5 Square-control bonus

6.6 Hanging pieces, defended pieces

6.6.1 SEE

6.6.2 SOMA

6.6.3 SUPER-SOMA

Part III

The Monte Carlo Revolution

Chapter 7

MCTS

Monte-carlo tree search (MCTS) is a relatively new approach to solving games that has become extremely effective at certain kinds of games, such as GO, where traditional $\alpha\beta$ -approaches have difficulties, both due to the high branching factor and to the challenge of constructing a good evaluation function.

7.1 UCB

Chapter 8

Extensions

8.1 AMAF

8.2 RAVE

8.3 Incorporating domain knowledge

Part IV

Other

Chapter 9

UI

In this chapter we

9.1 GGS

9.2 UCI

9.3 UGI

9.4 UGE

9.5 XBoard

9.6 Fritz

Chapter 10

Time management

In this chapter we

10.1 Time controls

There are many different time-controls that exists for different games. For example, there may be a fixed amount of time for each side for all moves made in the game, or there may be a set amount of time for each move made (where the unused time may be accumulated for later (with a possible cap) or may simply be lost.)

- Chess: In chess, there are a variety of popular time-controls.
- Go:
- Arimaa: from <http://arimaa.com/arimaa/learn/matchRules.html> and where the 2011 Arimaa Open Classic and World Championship is being played to 60s/5m/75/0/4h/4m for the Swiss preliminary and 90s/5m/75/0/6h/5m for the FDE final.

The time specification 60s/5m/75/0/4h/4m means, for example,

- 60 seconds per move
- 5 minutes in the starting reserve
- 75% of remaining move time gets added to reserve
- 4 hours maximum limit for the game
- 4 minutes maximum turn time

In other words: On each turn a player gets 60 seconds per move. If a move is not completed in this time then the reserve time may be used. There is a starting reserve of 5 minutes. If the reserve time is used up and the player has still not made a move then the player will lose on time. If the move is made in less than 60 seconds then 75% of the remaining move time is added to the reserve. A players turn may not exceed more than 4 minutes regardless of how much time is in reserve. If the game is not completed within 4 hours, it will be stopped and the winner determined by score.

- Computer Olympiad: in the ICGA Computer Olympiads, the most common time control is ...

10.1.1 Arimaa Time Controls Specification

The time control used for Arimaa is given by a 6-part specification M/R/P/L/G/T where

- M is the number of minutes:seconds per move; required
- R is the number of minutes:seconds in reserve; required
- P is the percent of unused move time that gets added to the reserve; optional defaults to 100

- L is the number of minutes:seconds to limit the reserve; 0 means no limit; optional; defaults to 0
- G is the number of hours:minutes after which time the game is halted and the winner is determined by score. G can also be specified as the maximum number of moves by ending with 't'; 0 means no limit; optional; defaults to 0
- T is the number of minutes:seconds within which a player must make the move; 0 means no limit; optional; defaults to 0

Different time units for any of the time control fields can be specified by adding one of the following letters after the numbers. In such cases the letter serves as the separator and : should not be used.

- s - seconds
- m - minutes
- h - hours
- d - days

For example: 24h5m10s/0/0/0/60d means 24 hours, 5 minutes and 10 seconds per move and the game must end after 60 days. Such a time control may be used in a postal type match.

- The game time parameter (G) can also be specified in terms of maximum number of turns each player can make by adding the letter t after the number.
- On each turn a player gets a fixed amount of time per move (M) and there may be some amount of time left in the reserve (R).
- If a player does not complete the move within the move time (M) then the time in reserve (R) is used. If there is no more time remaining in reserve and the player has not completed the move then the player automatically loses. Even if there is time left in the move or reserve, but the player has not made the move within the maximum time allowed for moves (T) then the player automatically loses.
- If a player completes the move in less than the time allowed for the move (M), then a percentage (P) of the remaining time is added to the players reserve. The result is rounded to the nearest second. This parameter is optional and if not specified, it is assumed to be 100%.
- An upper limit (L) can be given for the reserve so that the reserve does not exceed L when more time is added to the reserve. If the initial reserve already exceeds this limit then more time is not added to the reserve until it falls below this limit. The upper limit for the reserve is optional and if not given or set to 0 then it implies that there is no limit on how much time can be added to the reserve.
- For practical reasons a total game time (G) may be set. If the game is not finished within this allotted time then the game is halted and the winner is determined by scoring the game. This parameter is optional and if not given (or set to 0) it means there is no limit on the game time. Also instead of an upper limit for the total game time, an upper limit for the total number of turns each player can make may be specified by adding the letter 't' after the number. After both players have taken this many turns and the game is not finished the winner is determined by scoring the game.
- For games which use a time per move of less than 1 minute, both players are always given 1 minute of time to setup the initial position in the first move of the game. If the setup is not completed in 1 minute then the reserve time (R) is also used. The unused time from the setup move is not added to the reserve time.

Some examples:

- Example 1: 0/5 means 0 minutes per move with 5 minutes in reserve (per player). This is equivalent to G/5 in Chess; it means each player has a total of 5 minutes of time to play. If a player runs out of time before the game is over, the player loses. This is known as Blitz or "Sudden Death" time control in Chess.

- Example 2: 0:12/5 means 12 seconds per move with 5 minutes in reserve and all of the unused time from each move is added to the reserve time. It is similar to "5 12" in Chess which means "Game in 5 minutes with a 12 second increment". After each move 12 seconds is added to the remaining time. This is known as Incremental time control in Chess.
- Example 3: 3/0 means 3 minute per move and no reserve time, but 100 percent of the unused time for each move is added to the reserve. This guarantees that each player will make at least 40 moves in 2 hours. This is similar to the "40/2" Quota System time control used in Chess.
- Example 4: 0:30/5/100/3 means 30 seconds per move with 5 minutes in reserve and 100% of the unused time from each move is added to the reserve time. When the reserve already exceeds the limit, more time is not added to it. When the reserve falls below 3 minutes more time can be added to it, but the reserve is capped at 3 minutes.
- Example 5: 4/2/50/10/6 This means 4 minutes per move with a starting reserve of 2 minutes. After the move 50% of the time remaining for the move (rounded to the nearest second) is added to the reserve such that it does not exceed 10 minutes. There is a limit of 6 hours for the game after which time the game is halted and the winner is determined by score.
- Example 6: 4/4/100/4/6 This means 4 minutes per move and a starting reserve of 4 minutes. 100% percent of the unused move time gets added to the reserve such that it does not exceed 4 minutes. There is a time limit of 6 hour for the game after which the winner is determined by score.
- Example 7: 4/4/100/4/90t This is the same as above, but the game ends after both players have made 90 moves. Thus it ends after move 90 of silver is completed.
- Example 8: 4/4/100/4/90t/5 This is the same as above, but the players may not take more than 5 minutes for each turn even if there is still time remaining in reserve.

10.2 Estimation of remaining time

10.3 Control

10.4 Pondering

Pondering is where the computer uses the time during the opponent's move to think.

The standard approach is to assume that the opponent will make the move given as the response in the principal variation (PV) and to start "thinking" based on this assumption. As long as this correctly predicts the move 50% or more it will pay off well. (If the opponent makes a different move, we have to restart the search and the effort is lost.)

10.5 Example: Crafty

Here we describe the method for time-control used in Crafty (version xx.xxx).

10.6 Example: Gerbil

The open-source chess engine Gerbil supports a basic sudden-death time-control with or without an increment; a fixed time per move; and moves-per-period time-control.

It has variables **base** and **extra** in the engine.

For a fixed time t per move, it is quite simple: use all the given time **base** = t and there is no extra time **extra** = 0.

For sudden-death with no increment, it uses

```

case tctINCREMENT:
if (pcon->tc.tmIncr) {
TM tmTarget;

// An increment time control tries not to go below the increment
// plus 30 seconds, or 6 times the increment, whichever is more.
//
// If it finds itself below that, it will the increment minus
// 1/10 of the time by which it is under the target.
//
// If it is over that, it will use the increment plus 1/20 of
// the time by which it is over the target.
//
// So the program will drift slowly down to the target, but if
// it gets a ways below it, it will tend to come up quickly.
//
// There are a zillion different increment time controls, and
// sometimes I could end up with a stupid base time. I try hard
// to make sure the base time doesn't go below zero, which would
// be very bad since my time quantity is an unsigned value.
//
// I will try hard to use >= half of the increment. I will also
// make sure to not use more than half the remaining time (which
// is taken care of by the code at "lblSet".
//
tmTarget = pcon->tc.tmIncr + 30000;
if (tmTarget < 6 * pcon->tc.tmIncr)
tmTarget = 6 * pcon->tc.tmIncr;
if (pcon->ss.tmUs >= tmTarget)
tmBase = pcon->tc.tmIncr + (pcon->ss.tmUs - tmTarget) / 20;
else // The following statement can go negative on an unsigned.
tmBase = pcon->tc.tmIncr - (tmTarget - pcon->ss.tmUs) / 10;
Assert(sizeof(long) == sizeof(TM));
if ((long)tmBase < (long)pcon->tc.tmIncr / 2)
tmBase = pcon->tc.tmIncr / 2;
} else {
//
// A zero-increment time control will use 1/30 of the remaining
// time down to 10 minutes, 1/40 down to one minute, and 1/60
// after that.
//
if (pcon->ss.tmUs >= 600000) // 10 minutes.
tmBase = pcon->ss.tmUs / 30;
else if (pcon->ss.tmUs > 60000) // 1 minute.
tmBase = pcon->ss.tmUs / 40;
else
tmBase = pcon->ss.tmUs / 60;
}
// This code dummy-checks "tmBase", assigns "tmExtra", then assigns
// an end-time based upon "pcon->ss.tmStart".
//
// First I check to see if I'm scheduled to eat more than half of
// my remaining time. I put a ceiling at that amount, right off.
//

```

```

// Next, I'll assign some emergency time. This is 3x the base time,
// with a ceiling on base + extra of 1/2 the remaining time. I will
// also allocate no emergency time if I have < 30 seconds on the
// clock.
//
// So there will be plenty of "extra" time if there is time left,
// otherwise there is little or none.
//
lblSet: if (tmBase > pcon->ss.tmUs / 2)
tmBase = pcon->ss.tmUs / 2;
tmExtra = (pcon->ss.tmUs < 20000) ? 0 : tmBase * 3;
if (tmBase + tmExtra > pcon->ss.tmUs / 2)
tmExtra = pcon->ss.tmUs / 2 - tmBase;
pcon->ss.tmEnd = pcon->ss.tmStart + tmBase;
pcon->ss.tmExtra = tmExtra;
break;
case tctFIXED_TIME:
tmBase = pcon->tc.tmBase; // Use *all* the time.
pcon->ss.tmEnd = pcon->ss.tmStart + tmBase;
pcon->ss.tmExtra = 0; // No emergency time.
break;
case tctTOURNEY:
movMade = pcon->gc.ccm / 2 + pcon->gc.movStart - 1;
while (movMade >= pcon->tc.cMoves)
movMade -= pcon->tc.cMoves;
movLeft = pcon->tc.cMoves - movMade;
Assert(movLeft >= 1);
//
// This expression is kind of nasty. It is:
//
// Time / (3/4 x Moves + 3)
//
// The point of the "3" is to make sure there's some extra time pad
// at the beginning, and the point of the 3/4 is to make sure that
// earlier moves take a little longer per move.
//
// This expression has been reduced to ...
//
// (4 * Time) / (3 * Moves + 12)
//
// .. by multiplying the top and the bottom by 4. This helps out
// with the integer math.
//
tmBase = (4 * pcon->ss.tmUs) / (3 * movLeft + 12);
goto lblSet;
}

// This is called a whole bunch of times via the engine/interface protocol.
// Given what is known about the time control, the current time, and the
// search depth, this may set "ptcon->fTimeout", which will eventually finish
// the search.

// This routine has a built-in defense against cases where someone might try
// to force a move before the first ply has been fully considered. It will

```

```

// simply ignore such cases.

void VCheckTime(PCON pcon)
{
    TM tmNow;

    if (pcon->ss.plyDepth == 1) // Can't time out before a one-ply
    return; // search is finished.
    if (pcon->smode != smodeTHINK)
    return;
    if ((tmNow = TmNow()) >= pcon->ss.tmEnd) {
        //
        // Time has expired. Check for a fail-high or fail low and try to
        // add some time. Don't add any time if we're already dead lost and
        // just slowly collapsing, or if we are won and are probably failing
        // high repeatedly. The standard for "won" is a little greater than
        // for "lost".
        //
        if ((pcon->ss.prsa != prsaNORMAL) && (pcon->ss.val < valROOK) &&
            (pcon->ss.val >= -valMINOR))
            if (pcon->ss.tmExtra < 1000) {
                pcon->ss.tmEnd += pcon->ss.tmExtra;
                pcon->ss.tmExtra = 0;
            } else {
                pcon->ss.tmEnd += pcon->ss.tmExtra / 2;
                pcon->ss.tmExtra -= pcon->ss.tmExtra / 2;
            }
        if (tmNow >= pcon->ss.tmEnd)
            pcon->fTimeout = fTRUE;
    }
}

```

10.7 Example: Glaurung

As an interesting case-study, we examine the logic used by the open-source chess engine Glaurung. (From examining the source code.)

The basic timing method involves three variables to control the time spent on a search: **max_time** — the target time to spend on the search, **extra_time** — a possible amount of extra time to search in certain cases, and **absolute_max_time** — an absolute maximum, which will abort the search when exceeded. It supports sudden-death (with or without increment) or moves-per-period time-controls.

For sudden-death time-control with t remaining time (in milliseconds) and no increment, it sets:

$$\begin{aligned}\text{max_time} &= t/40 \\ \text{absolute_max_time} &= t/8\end{aligned}$$

If there is an increment of i milliseconds, then it uses

$$\begin{aligned}\text{max_time} &= t/30 + i \\ \text{absolute_max_time} &= \max(t/4, i - 100)\end{aligned}$$

For moves-per-period style of time-control, it sets

$$\begin{aligned}\text{max_time} &= t / \min(\text{moves_to_go}, 20) \\ \text{absolute_max_time} &= \min(4t/\text{moves_to_go}, t/3)\end{aligned}$$

unless there is only a single move left to go in which case it uses

$$\begin{aligned}\text{max_time} &= t/2 \\ \text{absolute_max_time} &= \min(t/2, t - 500)\end{aligned}$$

Glauring uses an iterative-deepening search strategy, and for the first 5 iterations, it sets `extra_time` to zero, otherwise (for the 6th and higher iterations) if the best move has changed in the last two iterations it will set

$$\text{extra_time} = \text{bmcc}_{it} \times (\text{max_time}/3) + \text{bmcc}_{it-1} \times (\text{max_time}/6)$$

where `bmccit` is the "best-move change-count" — the number of times that the best move has changed in the *it*'th iteration.

Glauring will stop a search early if there seems to be an "easy" move (a move which is clearly better than the rest). If it has done at least 5 iterations and:

```
// Stop search early if one move seems much better than the rest:
if(RSI->iteration >= 5 && RSI->easymove &&
  ((RSI->root_moves.moves[0].cumulative_nodes > (RSI->nodes*85)/100 &&
    get_time() - RSI->start_time > RSI->max_time/16) ||
  (RSI->root_moves.moves[0].cumulative_nodes > (RSI->nodes*99)/100 &&
    get_time()-RSI->start_time > RSI->max_time/32)) &&
  !RSI->infinite && RSI->thinking_status!=PONDERING) {
  break;
}
```

The search will be ended if most of `max_time` has been used at the end of an iteration (because there wouldn't be time to complete the next iteration in any case.) Specifically if the time used so far exceeds $(\text{max_time} + \text{extra_time}) \times 80/128$.

During the search, the timeout is

```
void check_for_timeout(void) {
  int t;
  static int last_info_time;

  t = get_time() - RSI->start_time;
  if(t < 1000) last_info_time = 0;
  else if(t - last_info_time >= 1000) {
    mutex_lock(IOLock);
    printf("info nodes " ll_u_format " nps " ll_u_format " time %d hashfull %d\n",
      RSI->nodes, (RSI->nodes*1000ULL)/((uint64) t), t, hashfull());
    mutex_unlock(IOLock);
    last_info_time = t;
    if(Options->currline) Threads[0].print_currline = true;
  }
  if(RSI->thinking_status != PONDERING && RSI->iteration >= 2 &&
    ((!RSI->infinite &&
      (t > RSI->absolute_max_time
        || (RSI->root_moves.current == 1 && t > RSI->max_time+RSI->extra_time)
        || (!RSI->fail_high && t > 6 * (RSI->max_time+RSI->extra_time)))) ||
      (RSI->node_limit && RSI->nodes >= RSI->node_limit) ||
      (RSI->exact_time && t >= RSI->exact_time)))
    RSI->thinking_status = ABORTED;
}
```

Chapter 11

Advanced Search

In this chapter we investigate more deeply the issue of searching,

11.1 Trappy min-max

This is a game-independent method described in [GR] for automatically playing in a way to set traps for the opponent — based on the idea that if you search more deeply than your opponent you can identify misleading lines of play (when the best-move to play at a shallower search depth is given a much lower evaluation at a deeper search depth). You may then choose a non-optimal move as long as it is not too much worse than the best move, but for which the obvious reply by the opponent is actually much worse ...TODO...

This kind of approach is best used against “fallible” opponents (human players, for example) or by brute-force deep-search based programs against shallower searching ... blah ...

11.2 Negascout

11.3 Null-window search

11.4 Aspiration search

11.5 PVS


11.6 MTD(f)

11.7 Conspiracy number

Conspiracy number search is a best-first search method. It is based on the idea of searching to a given level of confidence, measured in the number of nodes we’ve evaluated which would need to “conspire” in being wrongly evaluated in order to change our evaluation of the root. (We may also think of this as a measure of robustness — how many nodes could be systematically mis-evaluated without affecting our result.)

If we have a search tree, there will be a certain number of “leaf” nodes that we have reached in our search. (That is, nodes where we use the evaluation function to get an evaluation of the position.) Define the conspiracy number of *any* node in the tree to be the minimum number of *leaf* nodes that would need to change their value in order to change the value of that given node. In fact, for different evaluation values, there may be different numbers ... blah blah....

Standard example repeated *ad nauseum*:



where for node A we have

value	cn	nodes which have to change
1	2	(D or E) and (F or G)
2	1	F or G
3	0	-
4	1	E or F
5	1	E
6	2	(D and E) or (F and G)

11.7.1 Alpha-beta conspiracy search

In [?], McAllester and Yuret describe a variant of $\alpha\beta$ -search which is derived from conspiracy number search but is a depth-first rather than best-first search as in the original conspiracy numbers algorithm.

The idea incorporates two “depth” numbers d_{\max} and d_{\min} instead of a single one as in standard $\alpha\beta$ -search. It also includes an additional list of lists of static values of siblings of a given node and its ancestors.

We also define a number **delta**, the “singular margin”, as well as a real-valued, increasing function $S(k)$ such that $S(0) = 0$ and $\lim_{k \rightarrow \infty} S(k) = 1$ and which gets close to 1 before the typical branching factor b . (For example, we might take

$$S(k) = \begin{cases} 1 & \text{if } k \geq b \\ 1 - \left(\frac{b-k}{b}\right)^\gamma & \text{otherwise} \end{cases}$$

where b is the typical branching factor of the game we are analyzing. The parameter γ can be tuned from 1 to b to move continuously from conspiracy depth to classical ply depth.)

To start the search, one calls

```
ABC(root, -infinity, +infinity, d_max, d_min, [], [])
```

where

```
ABC(node, alpha, beta, d_max, d_min, max_options, min_options) {
  if (terminate?(node, alpha, beta, d_max, d_min, max_options, min_options)) {
    return the static value of node;
  }
  if (node is a max-node) {
    v := alpha;
    for each child c of node {
      sib_options = list of static values of siblings of c;
      next_max_options = [sib_options] ++ max_options;
      v = max(v, ABC(c, v, beta, d_max, d_min, next_max_options, min_options));
      if (v >= beta) { return v; }
    }
    return v;
  } else {
    ... dual of the max-node case ...
  }
}
```

and

```

terminate?(node, alpha, beta, d_max, d_min, max_options, min_options) {
    v = static value of node;
    if ((v >= beta) && (max_depth(beta, max_options) >= d_max)) {
        return true;
    }
    if ((v <= alpha) && (min_depth(alpha, min_options) >= d_min)) {
        return true;
    }
    if ((alpha < v < beta) && (max_depth(v, max_options) >= d_max) && (min_depth(v, min_options) >= d_min)) {
        return true;
    }
    return false;
}

```

and

```

max_depth(v, max_options) {
    return sum over all value lists L in max-options of S(k)
        where k is the number of elements of L > v-delta
}

```

and

```

min_depth(v, min_options) {
    return sum over all value lists L in min-options of S(k)
        where k is the number of elements of L < v+delta
}

```

11.8 Proof number

Implementation idea: to reduce space, assume pre-computed max number of nodes to search. Then we can pre-allocate huge array of elements. Children of node are allocated in adjacent sweep of elements from array - so parent just needs index of first child (and number of children). Note that this doesn't work well with strategies that free nodes etc.

Proof number search is a best-first search method which, as the name suggests, attempts to "prove" a result for a position. It does not rely on an evaluation function *per se*, but rather on a binary proposition or goal, for example, "Win for white" or "Not a loss for black". One can also imagine other propositions such as "White wins a queen". The search process then attempts to prove or disprove the proposition. In the discussion below, we will simply assume that the goal is "Win" for the max player for ease of exposition.

The *proof number* of a node is the number of nodes below it that need to be shown to be wins for the node itself to be a win and the *disproof number* is the number of its descendents that need to be proved losses for the node to be a loss. For a non-terminal leaf-node, we set both the proof and disproof numbers to 1. Thus, for a terminal node which is a win, we have a proof number of zero (0) and a disproof number of infinity (∞). Similarly, for a terminal node which is a loss, the proof number is ∞ and the disproof number is 0. For internal nodes, we can compute inductively their (dis)proof numbers: if a node X is a Max node, then the proof number $X.p$ and disproof number are given by

$$\begin{aligned}
 X.p &= \min_{C \in \text{Child}(X)} C.p \\
 X.d &= \sum_{C \in \text{child}(X)} C.d
 \end{aligned}$$

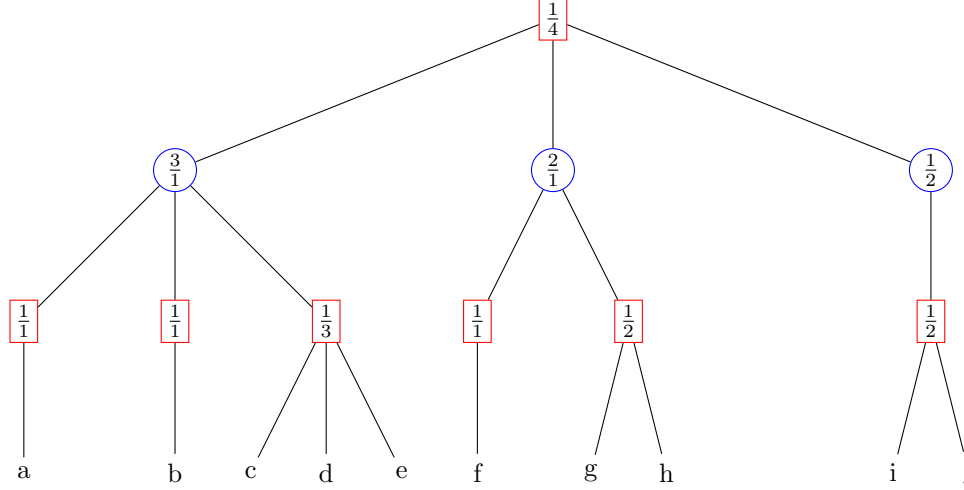
since to show a win we only need to find *one* winning move, and to prove a loss *every* move must be shown to lose. Similarly, for a Min node Y , we have

$$Y.p = \sum_{C \in \text{Child}(Y)} C.p$$

$$Y.d = \min_{C \in \text{child}(Y)} C.d$$

since our opponent needs to find only one move she can make that spoils our win to prove a loss, and for us to have a forced win, every possible move by the opponent must lead to a win.

Consider the following simple example:



We see that the proof number of the root node is 1 — this makes sense, since if either i or j are wins, then the rightmost move at the root node leads to a win. Similarly, the disproof number of the root node is 4 — this also makes sense, since for the root to be losing, each of the three possible moves must be losing, this requires at least 1 losing move for the first move (a or b) and at least 1 for the second (f) and both i and j must be losing.

Note that there are different sets of moves that would lead to a loss for the root node, for example if c, d, e, g, h, i, j are all losses, then the root is a loss, but that is not the *smallest* set of nodes which could disprove the root. The smallest possible disproof sets are a, f, i, j , or b, f, i, j . Similarly, if a, b, c are all wins, then the root is a win, but that set is not the smallest proving set.

...
The next step is to expand the “most-proving” node — this will be a node which is in both the smallest proving set and the smallest disproving set. In our case, either node i or node j are candidates. Thus one of them would be the next node expanded in the search. Its children would be expanded and the proof-numbers updated.

...

```

struct Node {
    mut internal : bool;
    mut children : [~Node];
    mut pn : int;
    mut dp : int;
}
fn determine_most_proving_node(V : &Node) -> &Node {
}

func determine_most_proving_node(V) {
    while (V is internal) {
        if (V is max_node) {

```

```

        V = leftmost child with equal proof number;
    } else {
        V = leftmost child with equal disproof number;
    }
}
return V;
}

func update(V) {
    if (V is max_node) {
        V.pn = minimum for S in children(V) of S.pn;
        V.dp = sum for S in children(V) of S.dp;
    } else {
        V.pn = sum for S in children(V) of S.pn;
        V.dp = minimum for S in children(V) of S.dp;
    }
    if (V has parent) {
        update(parent(V));
    }
}

```

Note that we can actually stop this update procedure when the values stop changing — and can begin the search for the most-proving node from this stopping point also!

```

func expand(V) {
    children(V) = generate_children(V);
    for all S in children(V) do {
        create_node(S);
        make_edge(V,S);
        value = evaluation(S);
        if (value=="Win") {
            s.pn = 0; s.dp = infinity;
        } else if (value=="Not win") {
            s.pn = infinity; s.dp = 0;
        } else {
            s.pn = s.dp = 1;
        }
    }
}

```

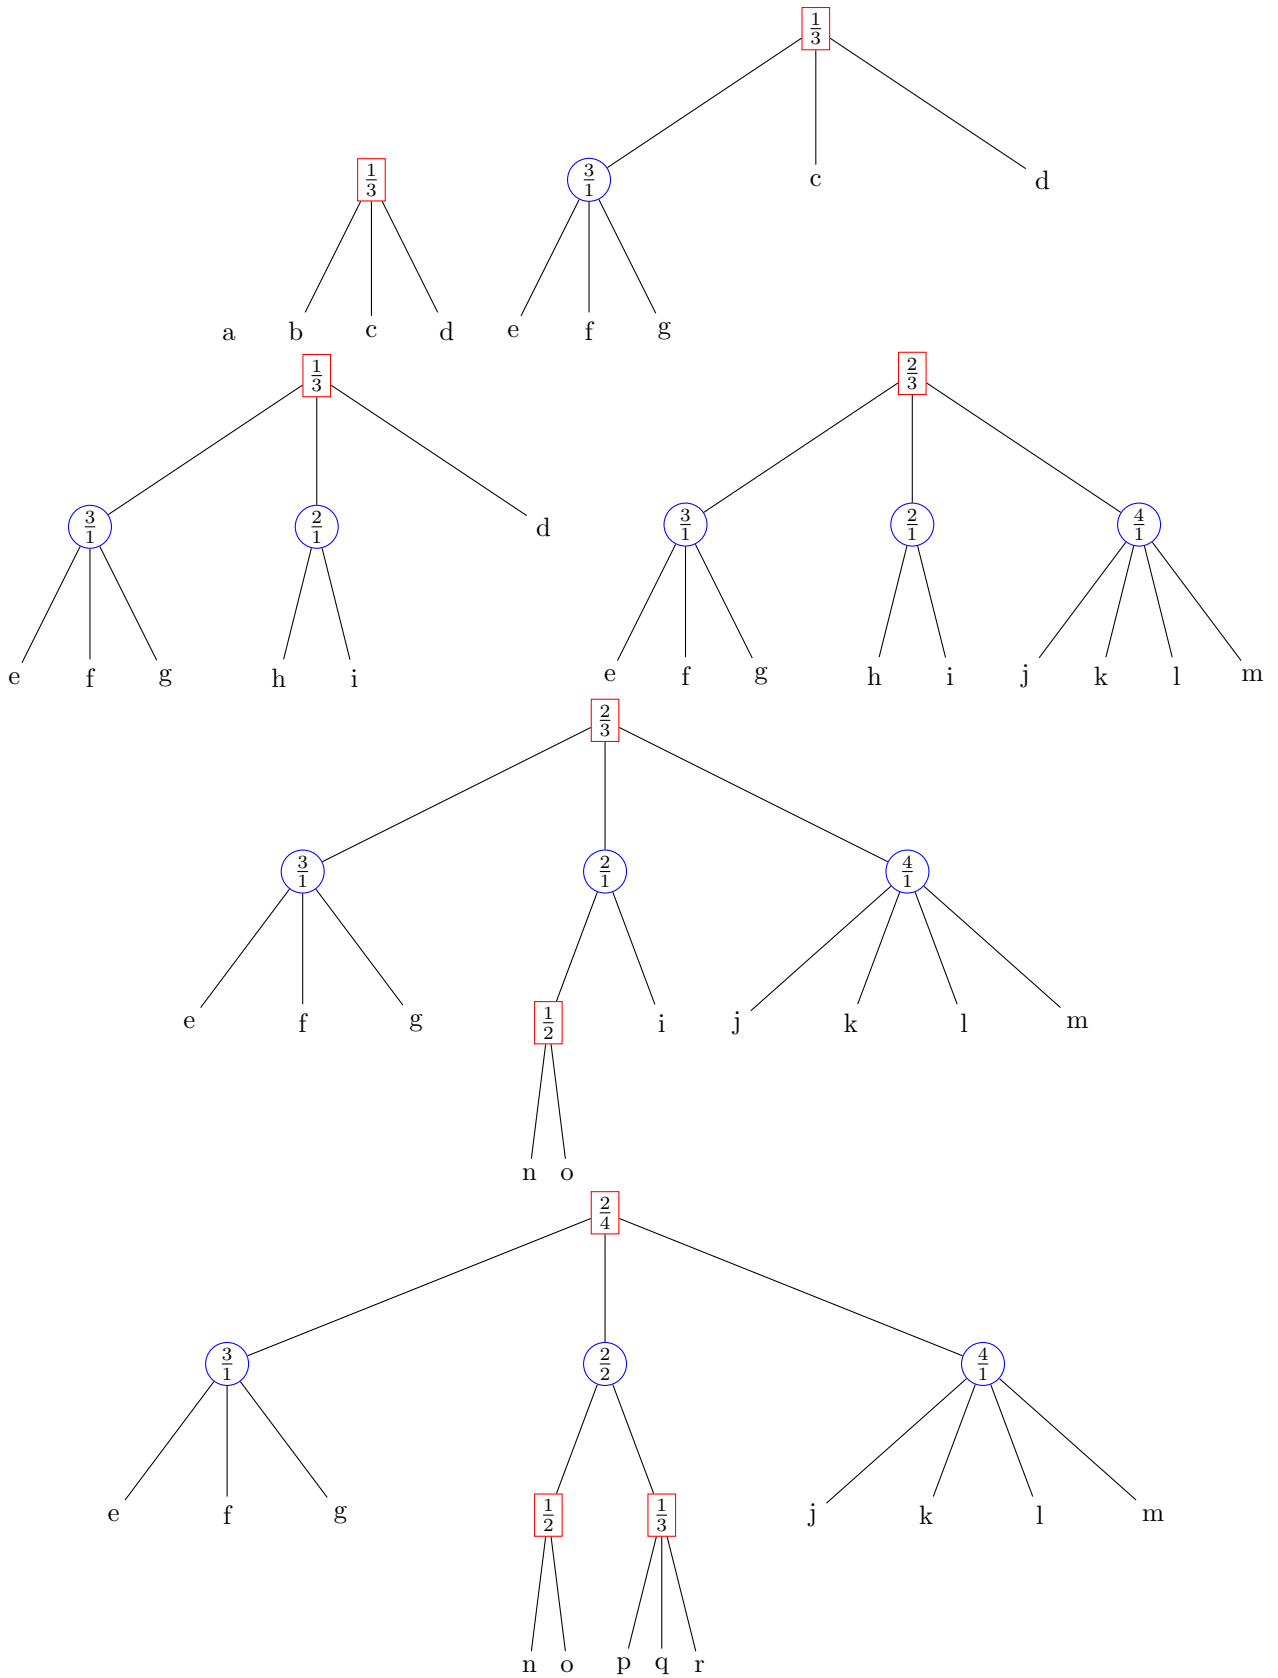
Could initialize leaves with other than 1/1 based on some heuristic (or even with the count of moves — it may be faster to just count moves than to generate them...)


```


func pn_search(position) {
    root = create_root(position);
    while (root.pn!=0 && root.dp!=0) {
        mpn = determine_most_proving_node(root);
        expand(mpn);
        update(mpn);
    }
    if (root.pn==0) return "Win";
    else return "Not win";
}

```

... step-by-step example of search & updated







...

11.8.1 Modifications and enhancements

Note that the proof-number algorithm (like the conspiracy numbers algorithm) works best on irregular and non-uniformly branching trees — if every node has exactly the same number of children and the game-end nodes are all at the same distance from root, then the algorithm degenerates into a breadth-first brute-force search.

Transposition tables... GHI problem...

11.9 Other heuristics

Chapter 12

Advanced Evaluation

In this chapter we discuss the evaluation of

12.1 TDleaf(λ)

One of the most difficult aspects of constructing an evaluation function can be “tuning” the weighting of different features.

Here we discuss a way for the computer to tune numeric parameters (such as piece values) automatically.

1. Let $J(x, w)$ be a heuristic evaluation function parametrized by $w = (w_1, \dots, w_k) \in \mathbf{R}^k$.
2. Let x_1, \dots, x_N be the positions that occurred in the game and let x'_i be the leaf node of the principal variation of the search starting at x_i .
3. Let $r(x_N) \in \{-1, 0, 1\}$ be the outcome of the game.
4. Define $r(x'_i, w) = \tanh(\beta \cdot J(x'_i, w))$, where $\beta = 0.255$ for knightcap, (where J is measured in units of 1 pawn).
5. For $i = 1, \dots, N-1$ compute $d_i = r(x'_{i+1}, w) - r(x'_i, w)$ and partial derivatives $\nabla r(x'_i, w) = (\dots \frac{\partial}{\partial w_j} r(x'_i, w) \dots)$.
6. If opponent rating is less than computer rating then set $d_i = 0$ if the move leading to x'_{i+1} was not predicted by the computer.
7. Update the parameters according to

$$w'_j = w_j + \alpha \sum_{i=1}^{N-1} \frac{\partial}{\partial w_j} r(x'_i, w) \left(\sum_{m=i}^{N-1} \lambda^{m-i} d_m \right)$$

where λ is around 0.3–0.7 if the evaluation function is already reliable and > 0.95 otherwise. Choose α so that updates are about 1/100 of a pawn.

Remarks on evaluating the partial-derivative of the evaluation function: in some cases, we may know the form of the evaluation and we can analytically derive the partial derivative. But suppose we cannot do this easily (or we simply can't be bothered to do it). A general approach to computing ...

$$\frac{\partial}{\partial z} f(\dots, z, \dots) \approx \frac{f(\dots, z+h, \dots) - f(\dots, z-h, \dots)}{2h}$$

for some small h (but not too small). This costs two computations of the evaluation function, but works without special knowledge of the evaluation function ...

12.1.1 TD(μ)

The TD(λ) algorithm suffers from the fact that it can learn to play badly from bad examples. The TD(μ) algorithm is a modification which is designed to allow learning to happen even with bad examples. This is especially useful in the case of self-play; when there is no good opponent to practice and learn from.

12.2 Neural networks

Neural networks provide a way

Neural networks have been used as components in various well-known and high-quality game-engines, including NeuroGammon by ...,

The down-side to neural networks is that executing them tends to be relatively slow, meaning that they are not generally amenable to being used in a deep-searching approach which requires evaluating millions of positions.

12.3 Genetic algorithms

An interesting, but ... Co-evolution ...

12.4 Pattern database / brains

Chapter 13

GUI

In this chapter we

Chapter 14

Parallel Processing

Parallel processing is a technique that is important for taking advantage of the processing power we have in current computers: from multi-core chips to massively-parallel GPUs to networked workstations. It is not straightforward to extend techniques such as $\alpha\beta$ -search to work efficiently in a parallel context. There have been a number of

The biggest issue with parallel extensions of $\alpha\beta$ -search is that the search is inherently sequential: the results of search of one part of the tree is used when searching another, namely the $\alpha\beta$ -bounds.

14.1 Youngest Brothers Wait

Chapter 15

Stuff

In this chapter we

15.1 Notation

15.1.1 Forsythe-Edwards Notation (FEN)

15.1.2 GBR code

15.1.3 PGN

15.2 Typesetting

15.3 Ratings

15.3.1 Elo

The Elo system (developed by Arpad Elo (?!)) is based on assigning numerical ratings to players under the statistical assumption that the probability that a player with rating A will beat a player with rating B is given by

$$\frac{10^{A/400}}{10^{A/400} + 10^{B/400}} = \frac{1}{1 + 10^{(B-A)/400}}$$

which means that a difference of 400 rating points corresponds to 10:1 odds in winning.

15.3.2 Glicko

15.4 Computer analysis

15.4.1 Blunder check

15.4.2 Auto-annotation

15.4.3 Combination finding

15.4.4 Automated problem construction

Chapter 16

Opening Book

In this chapter we discuss the usage and construction of an opening book. [Bur], [DL]

16.1 Basic format

16.2 Basic usage

16.2.1 Hash approach

One method (useful for small books with only positive entries) is to simply store all positions reachable from the book into the hash-table and when searching, if there is a move which reaches a position in the table, to take one of such moves (at random). This approach will break down when you want to store moves which should not be played, or if you want to associate probabilities or other more sophisticated ...

16.3 Automatic construction

16.4 Learning, updating

Chapter 17

Endgame Tablebase

In this chapter we discuss the usage and construction of an endgame tablebase.

Note that the approach of endgame tablebases is not applicable to every game. It works well in the cases where the game “simplifies” towards the end of the game. For example, if pieces tend to get captured as the game goes on and there are fewer and fewer pieces on the board (e.g. chess, checkers, mancala) then this approach works well. If the position gets more complicated as the game goes on (e.g. go, hex) or captured pieces are recycled (e.g. shogi) then it may have limited utility.

17.1 Building

The idea behind ...

The reason we use this backward approach of “un-moving” is to drastically reduce the amount of computation we need to do. Most positions in a given pass of the database will not be relevant, so rather than evaluating all possible moves from all possible positions in a database to see if they are a mate-in- n position, we quickly scan through to find the (relatively) few mate-in- $(n - 1)$ positions and only for those positions we do a (un-)move generation step. Though un-move-generation may be complex to code in a game, this can reduce computation time by orders of magnitude (including reducing significantly the number of random-access hits to the database being constructed — this is important for cache behaviour in in-memory databases, as well as disk-access speed in very large databases on disk.)

TODO: simulate behaviour for naive “forward-moving” vs “backward-moving” for some simple game (or congo for example? — reusing existing code? — can be for simple sub-game LE+L, for example) — compute number of move-generations, random-accesses to database, overall clock-time.

17.1.1 Thompson approach

X. Thompson [?] gave an early algorithm for “retrograde analysis of certain endgames”. It uses 5 active bit-files (for small enough endgames, they may fit entirely in memory and simply be bitmaps):

- W : known white-to-move-and-win
- B : known black-to-move-and-lose
- W_i : latest white-to-move-and-win-in- i -moves
- B_i : latest black-to-move-and-lose-in- i -moves
- J_i : temporary work-file

Then the algorithm proceeds as follows:

- (Setup) — Set B and B_0 to all positions with black to move but checkmated.
 — Set $W = 0$.

– Set $i = 0$.

(A) For each B_i position, generate predecessors — add them to W and all positions not already in W form the W_{i+1} positions.

(B) For each W_{i+1} position, generate predecessors to J_{i+1} .

(C) For each J_{i+1} position, generate successors — if all of them are in W then the position is added to B_{i+1} . Finally, add B_{i+1} to B .

(Loop) Increment i and iterate (A),(B),(C) until no changes to B and W

```
// Setup
W = 0;
B = B[0] = BlackCheckmated;
ChangedB = ChangedW = true;
for (i=0; ChangedB || ChangedW; ++i)
{
    // (A)
    J = Predecessors(B[i]);
    W[i+1] = J & ~W;
    ChangedW = NonEmpty(W[i+1]);
    W = W | W[i+1];
    // (B)
    J = Predecessors(W[i+1]);
    // (C)
    B[i+1] = ProveSuccessors(J, W);
    ChangedB = !Subset(B[i+1], B);
    B = B | B[i+1];
}
```

17.1.2 Wu-Beal approach

Here we describe the algorithm given in [?] for constructing endgame tablebases. It is designed for handling large databases, minimizing what must be kept in main memory (“random access”) ... fast sequential access disks ... Thus we have the two full databases W/B accessed sequentially, one working bitmap S accessed sequentially, and only a single working bitmap R which must have random-access.

```
// Set random-access bitmap R to predecessors of elements in sequential bitmap S
R Predecessors(S);

// load into S or R elements of W or B satisfying predicate P
S/R Load(W/B, P);

// Add elements of S to W or B with predicate P
W/B Add(S, P);

// Only keep elements of S with all successors in R
S = ProveSuccessors(S,R)

// W, B: full depth-to-win databases, accessed sequentially
// S : sequential bitmap
// R : random-access bitmap
void toplevel()
{
    do_initialize();
```

```

N = 0;                                     // depth to mate or conversion
while (!done_white || !done_black)
{
    if (!done_white)                       // last pass added new positions
    {
        S = Load(W, WIN_IN(N));           // S = WTM win_in_n
        R = Predecessors(S);               // R = BTM predecessors of S
        S = Load(B, UNKNOWN);             // S = BTM unknown
        S = S & R;                          // S = BTM may lose_in_n
        R = Load(W, WIN_IN(<=N));          // R = WTM win_in_(n or less)
        S = ProveSuccessors(S,R);          // S = BTM lose_in_n
        B = Add(S, LOSE_IN(N));            // B += S
        if (distance_to_mate)               // distance_to_mate?
            S = LOAD(B, LOSE_IN(N));        // S = BTM lose_in_n
        R = Predecessors(S);               // R = WTM maybe win_in_(n+1)
        S = Load(W, UNKNOWN);             // S = WTM unknown
        S = S & R;                          // S = WTM win_in_(n+1)
        W = Add(S, WIN_IN_N(n+1));         // W += S
    }
    if (!done_black)                       // done for BTM?
    {
        S = Load(B, WIN_IN(N));           // S = BTM win_in_n
        R = Predecessors(S);               // R = WTM predecessors of S
        S = Load(W, UNKNOWN);             // S = WTM unknown
        S = S & R;                          // S = WTM may lose_in_n
        R = Load(B, WIN_IN(<=N));          // R = BTM win_in_(n or less)
        S = ProveSuccessors(S,R);          // S = WTM lose_in_n
        W = Add(S, LOSE_IN(N));            // W += S
        if (distance_to_mate)               // distance_to_mate?
            S = LOAD(W, LOSE_IN(N));        // S = WTM lose_in_n
        R = Predecessors(S);               // R = BTM maybe win_in_(n+1)
        S = Load(B, UNKNOWN);             // S = BTM unknown
        S = S & R;                          // S = BTM win_in_(n+1)
        B = Add(S, WIN_IN_N(n+1));         // B += S
    }
}
}
}

```

17.2 Encoding: symmetries and numbering

An important technique for reducing the time and space complexity of the endgame tablebases is the use of symmetries to collapse “equivalent” positions into a single entry. A typical example would be the CHESS endgame of King vs. King + Queen: one can rotate or reflect the board without affecting the game-theoretic value of the positions, thus one could assume

Another technique to reduce the number of positions that needs to be considered is to take care to eliminate impossible positions. For example, in CHESS,

17.3 Compressing

Run-length encoding is a popular compression method that is well suited for sequential access to a table. It typically compresses significantly as many endgame tablebases consist of long runs of draws.

Random-access into the table is much slower as one must scan through sequentially to find the appropriate part of the table, though this can be sped up by the judicious use of indexes to point into the table.

17.4 Using

(internal nodes, tradeoffs in search, etc.)

The API for accessing the database needs to be designed with these symmetries in mind. And the search algorithm should be designed to interface with the endgame database in a transparent manner that allows ... Additionally, it is typically up to the ... Allow fast first check if database exists... then look up position

Part V

Specific Games

Chapter 18

Standard Models

There are a couple classes of standard approaches (or “models”) that are applicable to

18.1 $\alpha\beta$ -search

Our standard approach, which we will call the *$\alpha\beta$ -search model*, (applicable to most chess-like games) is as follows: negamax search with $\alpha\beta$ -pruning, iterative-deepening, transposition tables, quiescence search (where applicable), null-move pruning. [PVS search?] The static evaluation functions will depend heavily on the game in question but will be tuned using the TDLeaf(λ) algorithm.

Endgame tablebases will be built where applicable using retrograde analysis.

Opening books constructed using ...

The standard time-management approach will be used.

18.2 MCTS

Another standard approach, which has proven effective for some games, such as GO and HEX, is the *MCTS model*: a Monte-Carlo tree search based on UCT with modifications to incorporate domain-specific knowledge and with progressive widening.

18.3 Etc.

non-deterministic games

incomplete information

multi-player games

Arimaa...

Chapter 19

Congo

In this chapter we delve more deeply into the game of CONGO.

We will apply our standard approach of $\alpha\beta$ -search with iterative deepening, transposition table and quiescence extensions (based on captures, checks, and drowning.) For static evaluation we'll use material values with a piece-square bonus as well as some CONGO-specific features such as TODO

We'll use the TDLeaf(λ) self-learning approach to arrive at the appropriate weightings.

For end-games, we'll go ahead and build 3, 4, and 5-piece table-bases.

Opening books will be built by the self-play and deep-search approach.

— lion, pawn, super-pawn, elephant, monkey, alligator, giraffe

19.1 Endgame Tablebase

We will build end-game tablebases for all endings with up to 4 pieces (plus lions). First, let's estimate the number of positions in endgames of the different types. A few remarks are in order here: first, since the lions are restricted to their 3×3 dens, there are at most 81 positions for the pairs of lions to be in, and we will assume that the white lion is in the left-most 6 squares in his den, giving us only $6 \times 9 = 54$ different lion dispositions. The pawn cannot occur in the last row (since it will promote when it reaches the last line), giving only $7 \times 6 - 2 = 40$ possible positions (2 of the squares are occupied by lions). For other piece types, there are no restrictions except for the general restrictions of only one piece per square, giving $49 - n$ possible squares, where n is the number of occupied.

L+L		54	54
LP+L	54×40	2160	2214
LX+L (x=S/E/M/A/G)	$54 \times (49 - 2)$	12 690(2538 each)	14 904
LP+LP	$54 \times 40 \times (40 - 1)$	84240	99 144
LX+LP (x=S/E/M/A/G)	$54 \times (49 - 2) \times (40 - 1) \times 5$	494 910(98 982 each)	594 054
LX+LX (x=S/E/M/A/G)	$54 \times (49 - 2) \times (49 - 3) \times 5$		
LX+LY (x,y=S/E/M/A/G;x \neq y)	$54 \times (49 - 2) \times (49 - 3) \times 5 \times 4$		
LXP+L (x=S/E/M/A/G)	$54 \times (49 - 2) \times (40 - 1) \times 5$	494 910(98 982 each)	
LXY+L (x,y=S/E/M/A/G)	$54 \times (49 - 2) \times (49 - 3) \times 5 \times 5$	2 918 700(116 748 each)	

The simplest encoding gives for each particular man-combination:

L+L	9×9	81	81
LX+L (x=P/S/E/M/A/G)	$9 \times 9 \times 49 \times 6$	23 814(3969 each)	23 895
LXY+L (x,y=P/S/E/M/A/G)	$9 \times 9 \times 49^2 \times 6^2$	7 001 316(194 481 each)	7 025 211
LX+LY (x,y=P/S/E/M/A/G)	$9 \times 9 \times 49^2 \times 6^2$	7 001 316(194 481 each)	14 026 527
LXY+LZ (x,y,z=P/S/E/M/A/G)	$9 \times 9 \times 49^3 \times 6^3$	2 058 386 904(9 529 569 each)	
LXYZ+L (x,y,z=P/S/E/M/A/G)	$9 \times 9 \times 49^3 \times 6^3$	2 058 386 904(9 529 569 each)	

19.1.1 Another attempt at counting endgame positions (10-20-2011)

Note that we can choose either a lion *or* a piece to be restricted to, say, the left side of the board. To indicate which piece is restricted, we use underlining below. Note also that pawns cannot be in the first or last rows.

<u>LL</u>	6×9	54	54
<u>LXL</u> (x=ACEGSZ)	81×28	13,608(2,268 each)	
<u>LPL</u>	81×20	1,620	
<u>LXLY</u> (x,y=ACEGSZ)	$81 \times 28 \times 46$	2,190,888(104,328 each)	
<u>LXYL</u> (x,y=ACEGSZ; x _i y)	$81 \times 28 \times 46$	1,564,920(104,328 each)	
<u>LXXL</u> (x=ACEGSZ)	$54 \times \frac{47 \times 46}{2}$	350,244(58,374 each)	
<u>LXLP</u> (x=ACEGSZ)	$81 \times 28 \times 35$	476,280(79,380)	
<u>LXPL</u> (x=ACEGSZ)	$81 \times 28 \times 35$	476,280(79,380)	
<u>LPLP</u>	$81 \times 20 \times 34$	55,080	
<u>LPPL</u>	$81 \times 20 \times 34$	55,080	
<u>LXXLY</u> (x,y=ACEGSZ)	$54 \times \frac{47 \times 46}{2} \times 45$	94,565,880(2,626,830 each)	
<u>LXYLZ</u> (x,y,z=ACEGSZ; x _i y)	$81 \times 28 \times 46 \times 45$	422,528,400(4,694,760 each)	

Chapter 20

Mancala

In this chapter we delve more deeply into the sowing game of MANCALA. Now, there are many different variants of this game; indeed, one might rather call this a *family* of games, rather than a game. Examples include OWARI, MANCALA, KALAH, BAO, ...

20.1 Kalah

Kalah

20.2 Awari

20.3 Mancala

20.4 Dakon

20.5 Bao

Chapter 21

Chess

In this chapter we delve more deeply into the game of CHESS.

Chapter 22

Checkers

In this chapter we delve more deeply into the game of CHECKERS.

CHECKERS is an old game

CHECKERS is solved — that it is known to be a theoretical draw.

CHECKERS is

Chapter 23

Xiangqi

In this chapter we delve more deeply into the game of XIANGQI (or CHINESE CHESS).

Chapter 24

Shōgi

In this chapter we delve more deeply into the game of SHŌGI and some of its many variations.

24.1 Shōgi

the game of SHŌGI.

24.2 Dōbutsu shōgi

the game of DŌBUTSU SHŌGI.

24.3 Tori shōgi

the game of TORI SHŌGI.

24.4 Poppy shōgi

the game of POPPY SHŌGI.

24.5 Whale shōgi

the game of WHALE SHŌGI.

24.6 Chu shōgi

The game of CHU SHŌGI is a fascinating game that was ... china... 15th century... with ...

The biggest problem it will present is simply one of scale — the game is played on a 12×12 board and there are ... pieces on the board at the start of the game. This gives a large branching factor, implying that it will be difficult to search very deeply; but ...

24.7 Tenjiku shōgi

The game of TENJIKU SHŌGI is similar to CHU SHŌGI but is even larger: it is played on

Despite its large size, it has a highly tactical flavor

Chapter 25

Tafl

In this chapter we delve more deeply into the game of HNEFATAFL. This is an ancient game and not without interest. Due to its venerability there is a high degree of variability in rules and many variations and variants.

The specific variation we will cover here is ...

25.1 Evaluation

One component of an evaluation function we will use for this game is the number of consecutive moves it would take the king to win.

25.2 Mate search

This game seems ideal for approaches such as proof-number search or λ -search.

Chapter 26

Go

In this chapter we delve more deeply into the game of GO (or WEIQI).

Chapter 27

Havannah

In this chapter we delve more deeply into the game of HAVANNAH.

For this game, MCTS (with various modifications) has proven to be highly effective.

Chapter 28

Hex

In this chapter we delve more deeply into connection games, in particular the game of HEX.

Chapter 29

Hive

In this chapter we delve more deeply into the game of HIVE.

HIVE presents some new complications for programming. It has no set board — in principle it is played on an infinite board. That is, since the pieces are placed ...

Chapter 30

Amazons

In this chapter we delve more deeply into the game of AMAZONS.

Chapter 31

Dots-and-boxes

In this chapter we delve more deeply into the game of DOTS-AND-BOXES.

Chapter 32

Exotica

32.1 Backgammon

The game of BACKGAMMON represents new challenges, the biggest being that of randomness — usage of dice complicates matters.

32.2 Poker

The game of POKER represents many new challenges. In addition to the randomness of the shuffled deck, we also have imperfect information: we don't know what cards the other players are holding.

32.3 Kriegspiel, Stratego

The games KRIEGSPIEL and STRATEGO represents the challenges of

32.4 Wargames, Advanced Squad Leader

The game ADVANCED SQUAD LEADER (ASL)

Chapter 33

Combinatorial game theory

In this chapter we discuss

Chapter 34

Zillions of Games

Zillions of Games is a game package for Windows that uses a “universal gaming engine” technology that allows you to play nearly any abstract board game or puzzle in the world.

Chapter 35

International Computer Olympiada (ICGA)

Game	Tournaments	Programs
ABALONE	1	2
AMAZONS	9	12
AWARI	5	9
BACKGAMMON	7	12
BRIDGE	5	9
CHECKERS	2	8
CHES	60	245
CHINESE CHES	13	36
CLOBB	2	3
CONNECT6	4	17
CONNECT-FOUR	1	4
DOMINOES	1	3
DOTS AND BOXES	4	8
DRAUGHTS	9	21
GINRUMMY	1	2
GIPF	1	2
GO	25	59
GO-MOKU	4	16
HAVANNAH	1	2
HEX	6	9
KRIEGSPIEL	2	4
LINES OF ACTION	7	9
NINE MEN'S MORRIS	1	2
OCTI	1	2
OTHELLO	4	31
PHANTOM Go	3	4
POKER	1	2
POOL	3	6
QUBIC	2	4
RENU	4	9
SCRABBLE	4	8
SHOGI	9	10
SURAKARTA	2	3

35.1 2008 Computer Olympiad (call for participation)

Official games:

ABALONE, AMAZONS, ARIMAA, BACKGAMMON, BAO, BRIDGE, CHINESE CHESS, CLOBBER, COMPUTATIONAL POOL, CONNECT-6, DIPLOMACY, DOMINOES, DOTS-AND-BOXES, 10X10 DRAUGHTS, GIFP, GO, 9X9 GO, HEX, KRIEGSPIEL, LINES OF ACTION, OCTI, OTHELLO, PHANTOM GO, POKER, SCRABBLE, STRATEGO, SHOGI, SURAKARTA

Willing to host more if there were strong engines, such as:

ATAXX, DVONN, MEDIOCRITY, ONYX, TAMSK, TWIXT, ZÈRTZ

35.2 2009 Computer Olympiad (competed)

AMAZONS, CHESS, CHINESE CHESS, CONNECT6, DRAUGHTS, GO, GO (9X9), HEX, LINES OF ACTION, SHOGI, HAVANNAH, KRIEGSPIEL, PHANTOM GO

35.3 2010 Computer Olympiad

35.3.1 call for participation

AMAZONS, BACKGAMMON, BRIDGE, CHESS, COMPUTATIONAL POOL, CONNECT6, CHINESE CHESS, DOTS AND BOXES, INTERNATIONAL DRAUGHTS, 19X19 GO, 9X9 GO, HEX, HAVANNAH, LINES OF ACTION, SHŌGI, and, SURAKARTA

35.3.2 participated

Chapter 36

Game rules, sample games, comparisons

Game	Average game length	Average number of moves	Estimated State Space	Classifications
Checkers				
Chess				
Awari				
Shogi				Clas-
Go				
Amazons				
Shogi				
sification codes: Endgame: convergent/divergent/stationary Information: perfect/imperfect(common)/imperfect(private)				
Determinism: deterministic/stochastic				
CHESS, CHECKERS, CHINESE CHESS, SHOGI, CONGO, GO, HEX, DOTS-AND-BOXES, AMAZONS, MANCALA/AWARI, BACKGAMMON, CONNECT-4, PHUTBALL, TAFL/TABLUT, OTHELLO, KONANE, FOX-AND-GEESE, NINE-MAN-MORRIS, KHET, ABALONE, YAVALATH, PENTALATH, PATAGONIA, (Browne's games) GIPF, YINSH, TZAAR, ZÉRTZ, DVONN, PÜNCT, BRIDGE, POKER, SCRABBLE, MONOPOLY, KRIEGSPIEL, YAHTZEE, STRATEGO, settlers of cataan, carcasonne, puerto rico, yahtzee, parchisi, clue(do),				
<i>Encyclopedia of Chess Variants</i> : progressive chess, alapo, alice chess, atomic chess, avalanche chess, barrier chess I, beirut chess, berolina chess, betza's chess, bird chess I, bomb chess, caissa, carnivore chess, chad, checkless chess, chess football, chezz, chu shogi, columbia common chess, congo, diana, dragon chess, dragonfly, eternity's children, extinction chess, grand chess, hexagonal chess (glinski's), jetan, kamikaze chess I, koopa chess, loonybird, losing chess, makruk, microchess I/II/III/IV, mimic chess I/II, minishogi, missile chess, rifle chess, null chess, one-dimensional chess, parallel time-stream chess (?), parallel world chess, the pawn's game I, petty chess, ploy, rotary, plague chess, push chess, quest-chess, quickchess, refusal chess, compromise chess, sjakti, smess, spite chess, star-trek chess, three kings chess, tori shogi, tenjiku shogi, ultima, wa shogi, whale shogi, yari shogi				
[Hei00]				

36.1 Congo

CONGO is a chess-like game played on an 7×7 board with

36.2 Chess

CHESS is played on an 8×8 board with

The pieces' basic movement is given as follows

- King - moves one step in any direction

×	×	×
×	K	×
×	×	×

- Queen - slides any number of steps in any direction

↖	↑	↗
←	Q	→
↙	↓	↘

- Rook - slides any number of steps in an orthogonal direction

	↑	
←	R	→
	↓	

- Bishop - slides any number of steps in a diagonal direction

↖		↗
	B	
↙		↘

- Knight - jumps one orthogonal and one diagonal step

	×		×	
×				×
		N		
×				×
	×		×	

- Pawn - moves one step forward (without the right of capturing) or captures diagonally forward

<i>c</i>	<i>m</i>	<i>c</i>
	P	

36.3 Shogi

Glossary

PNS — see proof-number search

proof-number search — A method of searching

Index

10x10 DRAUGHTS, 60
19x19 GO, 60
9x9 GO, 60
ABALONE, 59–61
ADVANCED SQUAD LEADER, 56
AMAZONS, 54, 59–61
ARIMAA, 60
ATAXX, 60
AWARI, 59, 61
BACKGAMMON, 56, 59–61
BAO, 44, 60
BRIDGE, 59–61
CHECKERS, 46, 59, 61
CHESS, 17, 18, 39, 45, 59–61
CHINESE CHESS, 59
CHINESE CHESS, 47, 60, 61
CHU SHŌGI, 48
CHU SHOŌGI, 48
CLOBBER, 59, 60
COMPUTATIONAL POOL, 60
CONGO, 42, 61
CONNECT-4, 61
CONNECT-6, 60
CONNECT-FOUR, 59
CONNECT6, 59, 60
DŌBUTSU SHŌGI, 48
DIPLOMACY, 60
DOMINOES, 59, 60
DOTS AND BOXES, 59, 60
DOTS-AND-BOXES, 55, 60, 61
DRAUGHTS, 59, 60
DVONN, 60, 61
FOX-AND-GEESE, 61
GINRUMMY, 59
GIPF, 59–61
GO (9x9), 60
GO-MOKU, 59
GO, 20, 41, 50, 59–61
HAVANNAH, 51, 59, 60
HEX, 41, 52, 59–61
HIVE, 53
HNEFATAFL, 49
INTERNATIONAL DRAUGHTS, 60
KALAH, 44
KHET, 61

KONANE, 61
KRIEGSPIEL, 56, 59–61
LINES OF ACTION, 59, 60
MANCALA, 10, 44, 61
MEDIOCRITY, 60
MINI-CHESS, 10
MINI-SHOGI, 10
MONOPOLY, 61
NINE MEN’S MORRIS, 59
NINE-MAN-MORRIS, 61
OCTI, 59, 60
ONYX, 60
OTHELLO, 59–61
OWARI, 44
PÜNCT, 61
PATAGONIA, 61
PENTALATH, 61
PHANTOM GO, 59
PHANTOM GO, 60
PHUTBALL, 61
POKER, 56, 59–61
POOL, 59
POPPY SHŌGI, 48
QUBIC, 59
RENJU, 59
SCRABBLE, 59–61
SHŌGI, 18, 48, 60
SHOGI, 59–61
STRATEGO, 56, 60, 61
SURAKARTA, 59, 60
TABLUT, 61
TAFL, 61
TAMSK, 60
TENJIKU SHŌGI, 48
TIC-TAC-TOE, 9
TORI SHŌGI, 48
TWIXT, 60
TZAAR, 61
WEIQI, 50
WHALE SHŌGI, 48
XIANGQI, 47
YAHTZEE, 61
YAVALATH, 61
YINSH, 61
ZÉRTZ, 61

ZÈRTZ, 60

disproof number, 31

proof number, 31

Bibliography

- [Bur] Michael Buro, *Toward opening book learning*.
- [DL] Chrilly Donninger and Ulf Lorenz, *Innovative opening-book handling*.
- [GR] V. Scott Gordon and Ahmed Reda, *Trappy minimax - using iterative deepening to identify and set traps in two-player games*.
- [Hei00] Ernst Heinz, *Scalable search in computer chess*, Vieweg, 2000.